

FUSION TREES

Data structures II



Members:

Ali Muhammad Asad

Iqra Ahmed

Mohit Rai

Saira Junaid

Anoosha Hasan



PROJECT INTRODUCTION

The implementation of Fusion Trees

Background on Fusion trees

- The name Fusion Tree, is believed to be derived from from the "Cold Fusion Debacle", which was popular in the 90s, essentially a form of Fusion at Room temperature
- The fusion trees were developed, under the assumption that "We assume throughout this paper that the number of data items that are present never exceeds 2^w , the universe size. (To cope with a larger number of items, we could proceed by storing bucket pointers in our data structure, with equal items placed in common buckets.) ", as

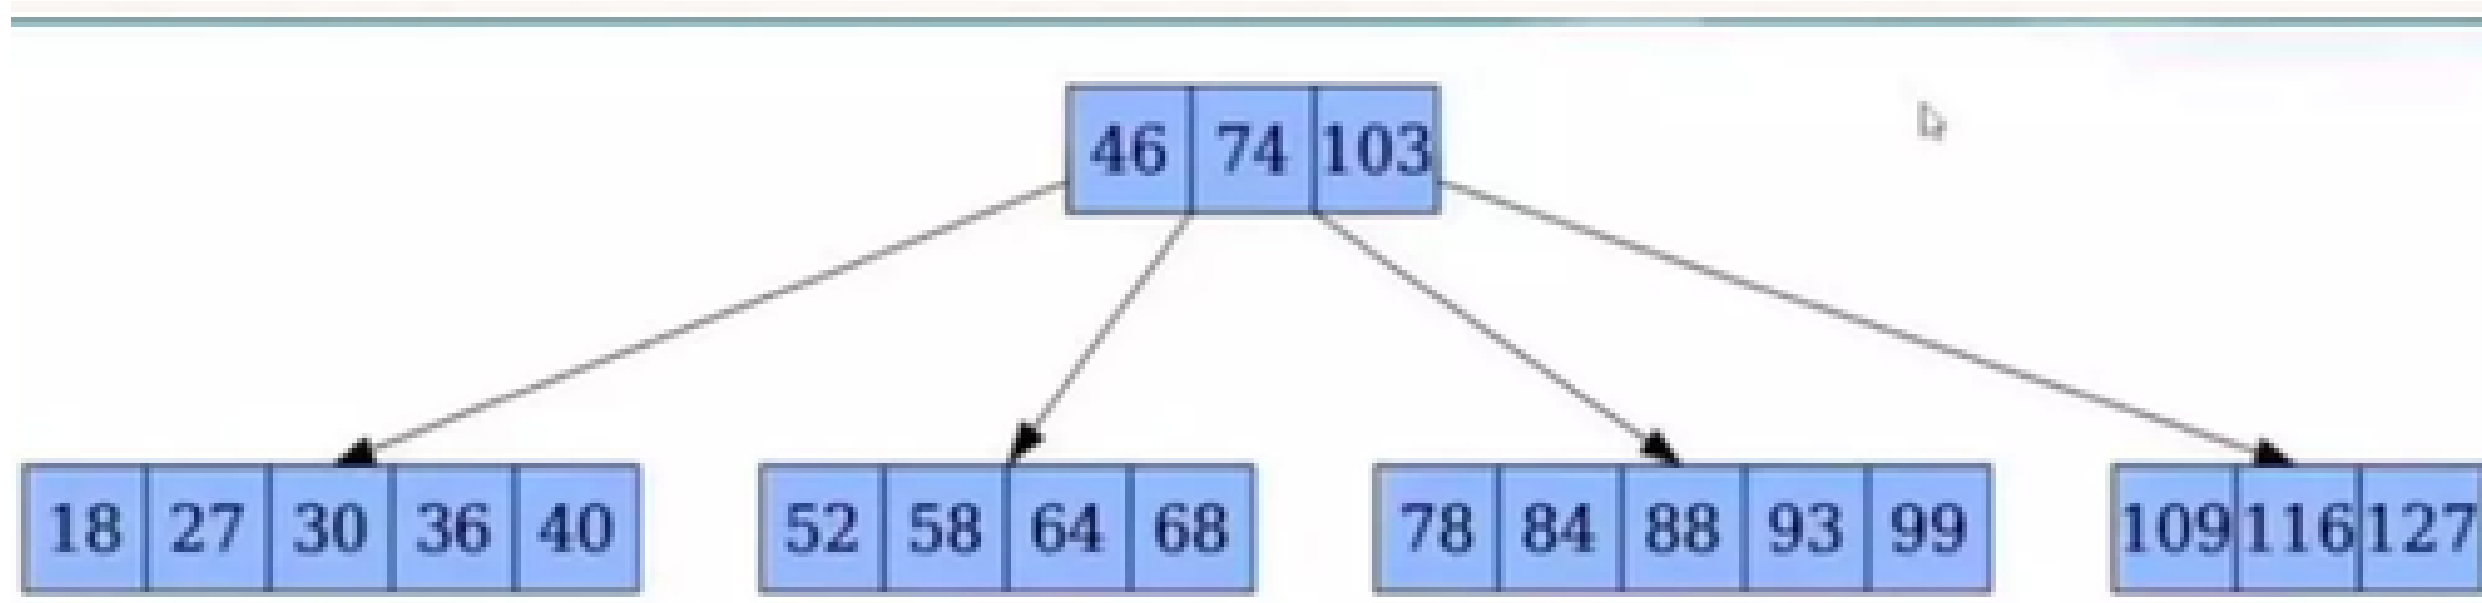
Fusion Trees

- A Fusion Tree is a modification of a B-Tree that implements an associative array on w -bit integers in a known universe size.
- An associative array is an abstract data type that stores a collection of pairs such as key-value pair.
- Fusion Trees are mostly used when our universe size is large, while providing linear - $O(n)$ - space complexity, and $O(\log n)$ search time complexity making it faster than a traditional self-balancing tree.

Fusion Trees Complexities

- Height: $O(\log w(n))$
 - Search, Predecessor, Successor: $O(\log w(n))$
 - Insert: $O(w^{2/3} \log w(n))$
 - Delete: $O(w^{2/3} \log w(n))$
 - Space: $O(n)$
-
- Where n is the count of values, w is word size ie, the size of the data type used to store the values.

Data structure of Fusion Tree



- Complexity of going through one fusion tree node becomes $O(1)$ thanks to "Parallel Comparison" of x with the keys in a node.
- Hence the search complexity of Fusion tree = $O(\log(t) n)$

The background features four decorative geometric patterns in the corners. Top-left: A series of parallel diagonal lines in a light blue-grey color, contained within a quarter-circle arc. Top-right: A cluster of quarter-circles in blue, yellow, red, and teal, some with a thin white grid pattern. Bottom-left: A cluster of quarter-circles in red, teal, and blue, some with a thin white grid pattern. Bottom-right: A series of parallel diagonal lines in a light blue-grey color, contained within a quarter-circle arc.

HOW DOES IT WORK?

How does it work?

- Essentially a B-Tree with a branching factor of $w^{1/5}$ giving it a height of $O(\log w(n))$ where 'w' is the machine word
- Our assumption is that universe size $\leq 2^w$
- Uses "Sketching" - fits keys all into one machine word allowing "Parallel Comparison" to be done.
- Helps achieve desired runtimes for updates and queries

The background features four decorative geometric patterns in the corners. The top-left corner has a series of parallel diagonal lines. The top-right corner contains a cluster of overlapping quarter-circles in blue, yellow, red, and teal, with a small blue circle containing a white 'a' nearby. The bottom-left corner shows a cluster of overlapping quarter-circles in red, teal, and blue. The bottom-right corner features a large, faint arc with several parallel diagonal lines inside it.

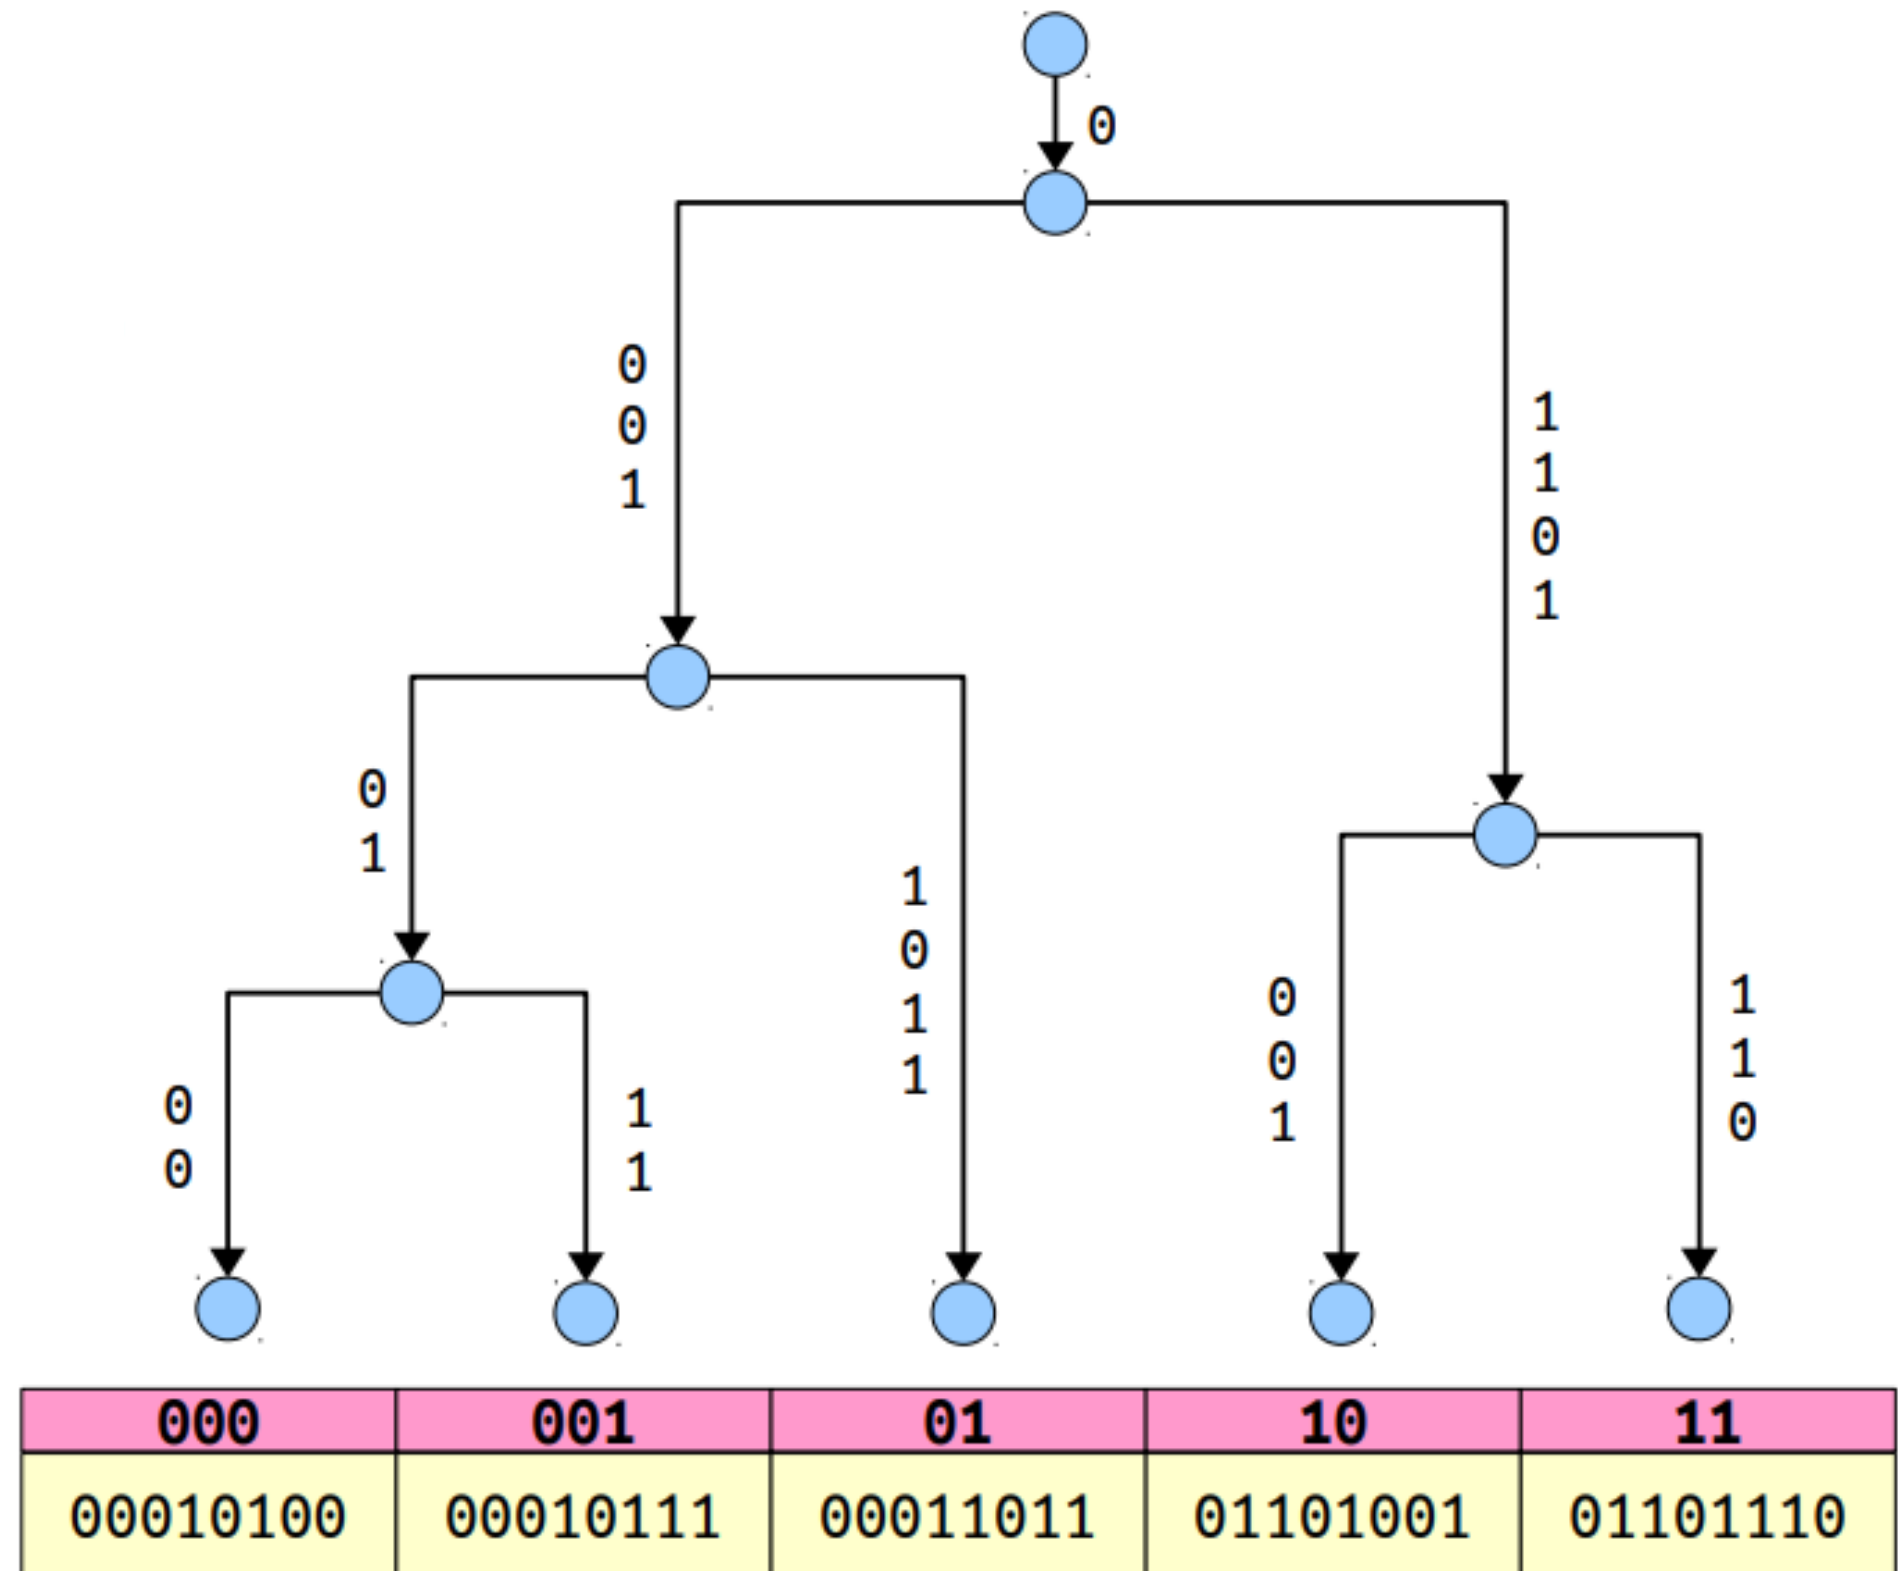
SKETCHING

Sketching

- Each w -bit key at a node containing k -keys is compressed
- Creates a path-like structure in a full binary tree of height w
- Max k -keys, so maximum $k - 1$ branching points (where two keys differ)
- Preserves order of the keys, that is $\text{sketch}(x) < \text{sketch}(y)$ for any two keys in a node where $x < y$.

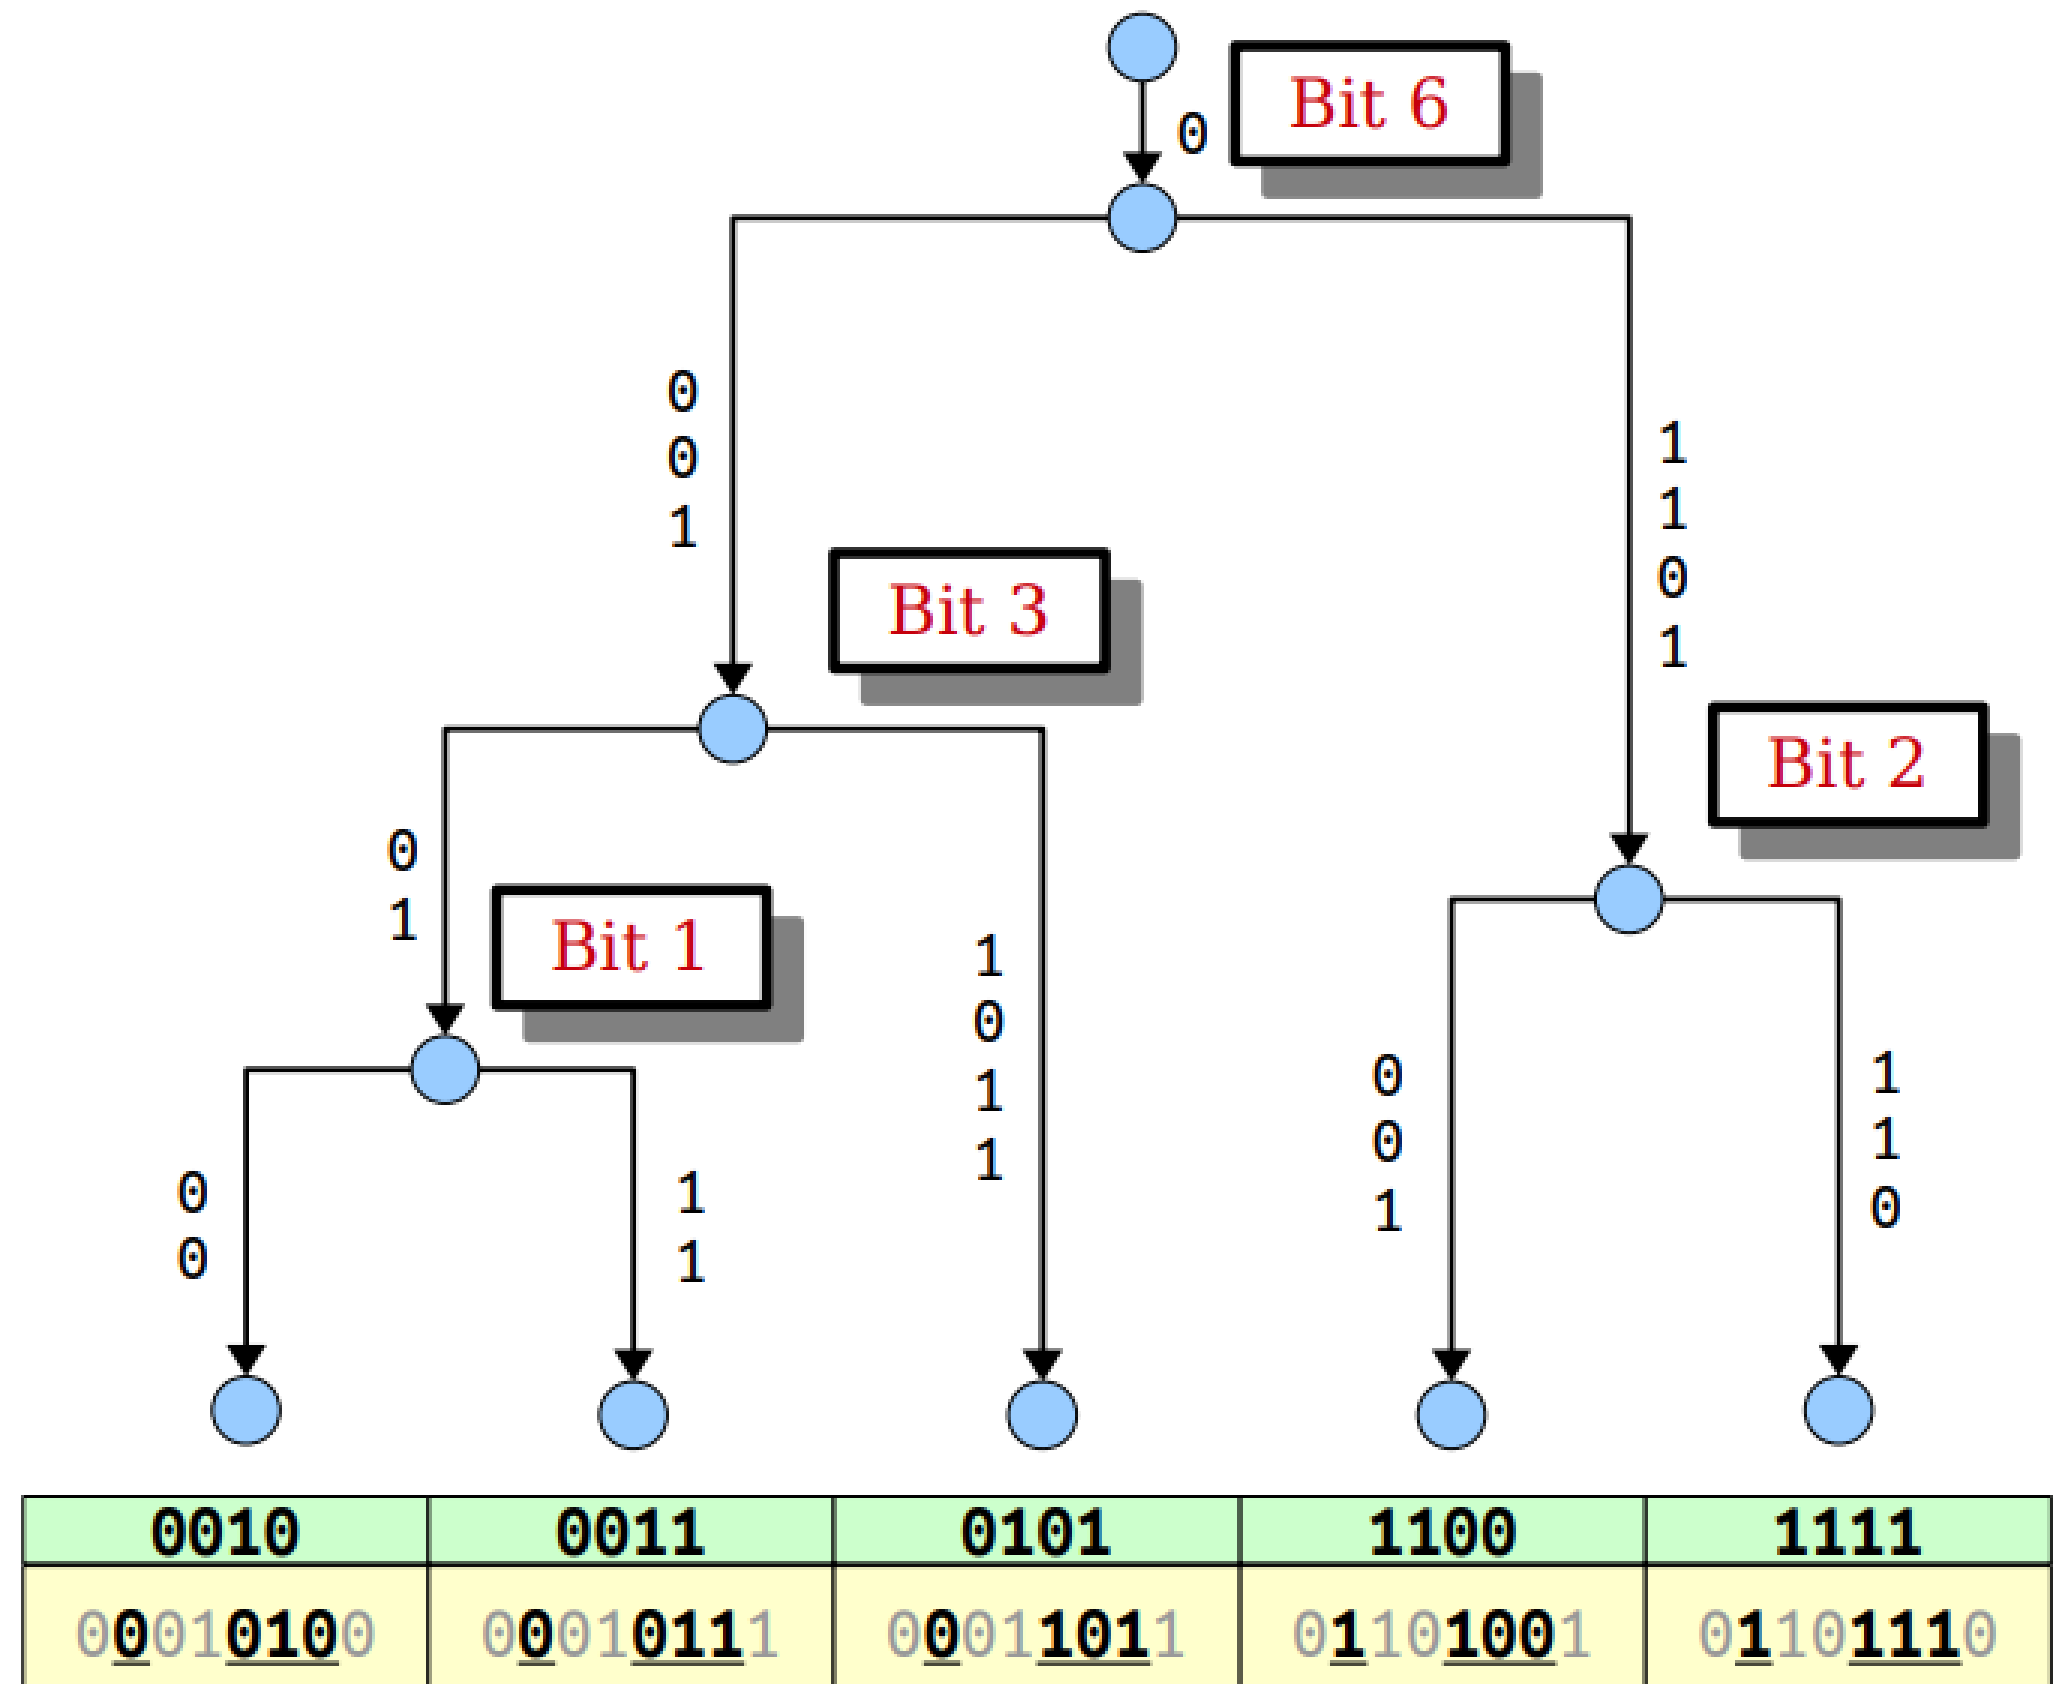
Sketching - Aims

- Ultimately we are interested in compressing our numbers so that they fit in a machine word.
- There at most w^ϵ (epsilon) bits in each of these new numbers - that's really promising!



Sketching - Patricia Code


- A bit index i is interesting if there is a branching node in the trie at that bit index
- The *Patricia Code* of an integer is the bitstring consisting of just the interesting bits in that number



Bringing it Together

$\text{sketch}(21) = \text{sketch}(10101) = 10\underline{1}01 = 01$
 $\text{sketch}(17) = \text{sketch}(10001) = 10\underline{0}01 = 00$
 $\text{sketch}(23) = \text{sketch}(10111) = 10\underline{1}11 = 11$

- Not possible to compare $O(k)$ keys of a node in constant time
- Sketch compresses the keys for faster comparison of keys within a node
- Allows constant time for predecessor queries





PARALLEL COMPARISON

Parallel Comparison

- Finds the position of a value within a set of keys in a node in constant time
- Provides efficient storage, and retrieval of large sets of items
- Multiple keys are compared simultaneously at each level
- Uses bitwise operations for comparison of the binary representation of keys given by Sketching

Parallel Comparison

	01101110	00101110	01111000	01001101	00101111	00001101	01110111	01100001
+	00011010	01000101	00010100	00100000	01010000	00100010	01000100	00001000
<hr/>								
	10001000	01110011	10001100	01101101	01111111	00101111	10111011	01101001

- Eight pairs of 7-bit numbers are compared using a single 64-bit addition
- The MSBs of each are summed to evaluate the appropriate positions


The background features four decorative geometric patterns in the corners. The top-left corner has a series of parallel diagonal lines. The top-right corner contains a cluster of overlapping quarter-circles in blue, yellow, red, and teal, with a small blue circle containing a white 'a' nearby. The bottom-left corner shows a cluster of overlapping quarter-circles in red, teal, and blue. The bottom-right corner features a large, faint arc with several parallel diagonal lines extending from its base.

APPLICATION

A	B	C	D	E
		Database clients of Jolly Da		
No	Customer	Type	Country	City
1	Intersection	com.network	USA	New York
2	Magnet	com.network	USA	New York
3	Perspective korp.	warehouse	Belarus	Minsk
4	Driveway	enterprise	USA	New York
5	near	enterprise	USA	Los Angeles
6	Nori	warehouse	Japan	Tokyo
7	Nevsky comp.	com.network	Russia	Moscow
8	Perspective korp.	enterprise	Belarus	Minsk
9	in touch	warehouse	USA	San Francisco
10	Nardis	com.network	Japan	Tokyo


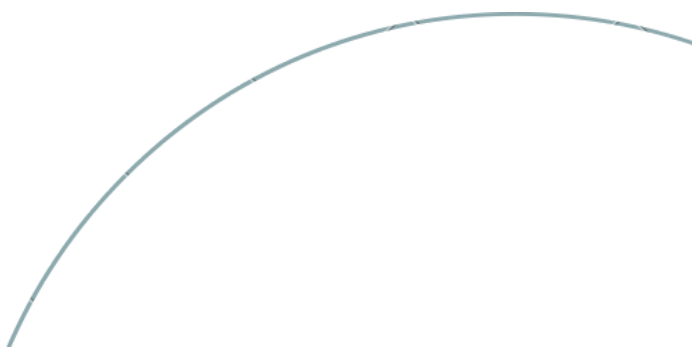
- Databases Management Systems for faster insertion, deletion, and lookup
- Indexes for search engines involving numerical data

- Computational Geometry
- Machine Learning
- Geometric Problems like nearest neighbor queries in higher dimensions




The background features decorative geometric patterns in the corners. The top-left and bottom-right corners contain light blue diagonal line patterns. The top-right and bottom-left corners feature clusters of overlapping quarter-circles in teal, yellow, and red. The central text is in a bold, blue, sans-serif font.

LIMITATIONS AND COMPLEXITIES

- 
- **Design: a unique combination of ideas and techniques from various mathematical fields - challenging implementation and understanding**
 - **Bit manipulation and bitwise operations for lookup - predecessor and successor search**
 - **Less efficient space complexity**
 - **Complex insertion and deletion**
 - **Non-intuitive design**
 - **Limited practicality**
- 

REFERENCES

- [1] Fredman, M. L.; Willard, D. E. (1990), "BLASTING Through the Information Theoretic Barrier with FUSION TREES", Proceedings of the Twenty-Second Annual ACM Symposium on Theory of Computing (STOC '90), New York, NY, USA: ACM, pp. 1–7, doi:10.1145/100216.100217, ISBN 0-89791-361-2, S2CID 16367160.
- [2] https://www.youtube.com/watch?v=xSGorVW8j6Q&ab_channel=MITOpenCourseWare
- [3] <https://youtu.be/QXASOE5scoM>
- [4] <https://youtu.be/Jec59qLhkGU>
- [5] <https://www.slideshare.net/RohithND/merkle-trees-and-fusion-trees>
- [6] <https://stackoverflow.com/questions/3878320/understanding-fusion-trees>



The background features several decorative geometric patterns. In the top-left corner, there are thin, parallel diagonal lines. In the top-right corner, there is a cluster of overlapping semi-circles in blue, teal, orange, and red. In the bottom-left corner, there is another cluster of overlapping semi-circles in blue, teal, orange, and red, with a small blue semi-circle containing a white lowercase 'v' at the top. In the bottom-right corner, there is a large, faint, light blue arc and some thin diagonal lines.

THANK YOU